



Deploy Kubernetes 1.9 from scratch on VMware vSphere

29 December 2017

This lab will go through the different steps needed to configure an HA Kubernetes cluster on VMware vSphere manually. The various communications between the Kubernetes components will be secured with TLS. If you are used to deploy Kubernetes with tools like kubeadm but would like to understand a bit more what is going on under the hood, this tutorial is for you.



(http://blog.inkubate.io)

However, it is good to have an understanding of the architecture of each Kubernetes node as described in the [Kubernetes components documentation](https://kubernetes.io/docs/concepts/overview/components/) (<https://kubernetes.io/docs/concepts/overview/components/>).

This article is inspired by the awesome tutorial [Kubernetes the hard way](https://github.com/kelseyhightower/kubernetes-the-hard-way) (<https://github.com/kelseyhightower/kubernetes-the-hard-way>), which explains how to deploy a Kubernetes cluster from scratch on Google Cloud Platform. Some of the steps are a bit different on VMware vSphere. For example, VMware vSphere doesn't provide a native load balancer and we will have to install and configure one ourselves.



kubernetes



(<http://blog.inkubate.io>)

Prerequisites

Deploy Kubernetes 1.9 from scratch on VMware vSphere

9%

For this lab, we will use a standard Ubuntu 16.04 installation as a base image for the seven virtual machines needed. The virtual machines will all be configured on the same virtual network 10.10.40.0/24 and this network needs to have access to the Internet.

The first machine needed is the machine on which the HAProxy load balancer will be installed. We will assign the IP 10.10.40.63 to this machine.

We also need three Kubernetes master nodes. These virtual machines will have the IPs 10.10.40.60, 10.10.40.61, and 10.10.40.62.

Finally, we will also have three Kubernetes worker nodes with the IPs 10.10.40.70, 10.10.40.71, and 10.10.40.72.

We also need an IP range for the pods. This range will be 10.20.0.0/16, but it is only internal to Kubernetes and doesn't need to be configured on VMware vSphere.

I will use my Linux desktop as a client machine to generate all the necessary certificates, but also to manage the Kubernetes cluster. If you don't have a Linux desktop, you can use the HAProxy virtual machine to do the same thing.

Installation of the client tools



We will need two tools on the client machine. The Cloud Flare SSL tool to generate the different certificates, and the Kubernetes client, kubectl, to manage the Kubernetes cluster.

(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words, 23 min read

9%

Installation of cfssl

1- Download the binaries.

```
$ wget https://pkg.cfssl.org/R1.2/cfssl_linux-amd64
```

```
$ wget https://pkg.cfssl.org/R1.2/cfssljson_linux-amd64
```

2- Add the execution permission to the binaries.

```
$ chmod +x cfssl*
```

3- Move the binaries to /usr/local/bin.

```
$ sudo mv cfssl_linux-amd64 /usr/local/bin/cfssl
```

```
$ sudo mv cfssljson_linux-amd64 /usr/local/bin/cfssljson
```

4- Verify the installation.

```
$ cfssl version
```

Installation of kubectl





- Download the binary
Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```

2- Add the execution permission to the binary.

```
$ chmod +x kubectl
```

3- Move the binary to /usr/local/bin.

```
$ sudo mv kubectl /usr/local/bin
```

4- Verify the installation.

```
$ kubectl version
```

Installation of the HAProxy load balancer

As we will deploy three Kubernetes master nodes, we need to deploy an HAProxy load balancer in front of them to distribute the traffic.

1- SSH to the 10.10.40.63 Ubuntu virtual machine.



2- Update the virtual machine. Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

```
$ sudo apt-get update
```

```
$ sudo apt-get upgrade
```

3- Install HAProxy.

```
$ sudo apt-get install haproxy
```

4- Configure HAProxy to load balance the traffic between the three Kubernetes master nodes.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

```
...
global
...
default
...
frontend kubernetes
    bind          10.10.40.63:6443
    option        tcplog
    mode          tcp
    default_backend kubernetes-master-nodes

backend kubernetes-master-nodes
    mode          tcp
    balance        roundrobin
    option        tcp-check
    server k8s-master-0 10.10.40.60:6443 check fall 3 rise 2
    server k8s-master-1 10.10.40.61:6443 check fall 3 rise 2
    server k8s-master-2 10.10.40.62:6443 check fall 3 rise 2
```

5- Restart HAProxy.

```
$ sudo systemctl restart haproxy
```

Generate the TLS certificates

These steps can be done on your Linux desktop if you have one or on the HAProxy virtual machine depending on where you installed the cfssl tool.

Create a certificate authority

1- Create the certificate authority configuration file.

```
$ vim ca-config.json
{
  "signing": {
    "default": {
      "expiry": "8760h"
    },
    "profiles": {
      "kubernetes": {
        "usages": ["signing", "key encipherment", "server auth", "client auth"]
        "expiry": "8760h"
      }
    }
  }
}
```

2- Create the certificate authority signing request configuration file.



(<http://blog.inkubate.io>)

\$ vi DeployKubernetes 1.9 from scratch on VMware vSphere

{ 6116 words - 23 min read.

```
{
  "CN": "Kubernetes",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "Kubernetes",
      "OU": "CA",
      "ST": "Cork Co."
    }
  ]
}
```

9%

3- Generate the certificate authority certificate and private key.

```
$ cfssl gencert -initca ca-csr.json | cfssljson -bare ca
```

4- Verify that the ca-key.pem and the ca.pem were generated.

```
$ ls -la
```

Generate the Kubernetes certificates



Admin client certificate

This certificate will be used to connect to the Kubernetes cluster as an administrator.

1- Create the certificate signing request configuration file.

```
$ vim admin-csr.json
{
  "CN": "admin",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "system:masters",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```

2- Generate the certificate and the private key.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ cfssl gencert \
  -ca=ca.pem \
  -ca-key=ca-key.pem \
  -config=ca-config.json \
  -profile=kubernetes admin-csr.json | \
  cfssljson -bare admin
```

9%

3- Verify that the admin-key.pem and the admin.pem were generated.

```
$ ls -la
```

Generate the Kubelet client certificates

We will deploy a kubelet on each Kubernetes worker node. Each of these kubelets will need a certificate to join the Kubernetes cluster.

1- Create a certificate signing request configuration file for the 10.10.40.70 worker node.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
{
  "CN": "system:node:10.10.40.70",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "system:nodes",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```

9%

2- Generate the certificate and the private key.

```
$ cfssl gencert \
-ca=ca.pem \
-ca-key=ca-key.pem \
-config=ca-config.json \
-hostname=10.10.40.70,10.10.40.70 \
-profile=kubernetes 10.10.40.70-csr.json | \
cfssljson -bare 10.10.40.70
```

3- Create a certificate signing request configuration file for the 10.10.40.71 worker node.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
{
  "CN": "system:node:10.10.40.71",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "system:nodes",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```

9%

4- Generate the certificate and the private key.

```
$ cfssl gencert \
-ca=ca.pem \
-ca-key=ca-key.pem \
-config=ca-config.json \
-hostname=10.10.40.71,10.10.40.71 \
-profile=kubernetes 10.10.40.71-csr.json | \
cfssljson -bare 10.10.40.71
```

5- Create a certificate signing request configuration file for the 10.10.40.72.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
{
  "CN": "system:node:10.10.40.72",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "system:nodes",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```

9%

6- Generate the certificate and the private key.

```
$ cfssl gencert \
-ca=ca.pem \
-ca-key=ca-key.pem \
-config=ca-config.json \
-hostname=10.10.40.72,10.10.40.72 \
-profile=kubernetes 10.10.40.72-csr.json | \
cfssljson -bare 10.10.40.72
```



Deploy Kubernetes 1.9 from scratch on VMware vSphere

7- Verify that the files 10.10.40.70-key.pem, 10.10.40.70.pem, 10.10.40.71-key.pem, 10.10.40.71.pem, 10.10.40.72-key.pem, and 10.10.40.72.pem were generated.

(<http://blog.inkubate.io>)

9%

```
$ ls -la
```

Generate the kube-proxy client certificate

Another component that we will install on all our worker nodes is the kube-proxy. The kube-proxy also needs a client certificate.

1- Create the certificate signing request configuration file.

```
$ vim kube-proxy-csr.json
{
  "CN": "system:kube-proxy",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "system:node-proxier",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```



2- Generate the kube-proxy certificate and private key.

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

```
$ cfssl gencert \  
-ca=ca.pem -ca-key=ca-key.pem \  
-config=ca-config.json \  
-profile=kubernetes kube-proxy-csr.json | \  
cfssljson -bare kube-proxy
```

3- Verify that the kube-proxy-key.pem and the kube-proxy.pem files were generated.

```
$ ls -la
```

Generate the API server certificate

The last certificate we need to generate is the API server certificate. The API component will be installed on each of our Kubernetes master nodes.

1- Create the certificate signing request configuration file.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

{ 6116 words - 23 min read.

```
{
  "CN": "kubernetes",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "IE",
      "L": "Cork",
      "O": "Kubernetes",
      "OU": "Kubernetes",
      "ST": "Cork Co."
    }
  ]
}
```

9%

2- Generate the API server certificate and private key.

```
$ cfssl gencert \
-ca=ca.pem \
-ca-key=ca-key.pem \
-config=ca-config.json \
-hostname=10.32.0.1,10.10.40.60,10.10.40.61,10.10.40.62,10.10.40.63,127.0.0.1,k
-profile=kubernetes kubernetes-csr.json | \
cfssljson -bare kubernetes
```

3- Verify that the kubernetes-key.pem and the kubernetes.pem file were generated.





(<http://blog.inkubate.io>)

Copy the certificates to the nodes

Worker nodes

1- Copy the certificates to the 10.10.40.70 worker node.

```
$ scp ca.pem 10.10.40.70-key.pem 10.10.40.70.pem 10.10.40.70:~
```

2- Copy the certificates to the 10.10.40.71 worker node.

```
$ scp ca.pem 10.10.40.71-key.pem 10.10.40.71.pem 10.10.40.71:~
```

3- Copy the certificates to the 10.10.40.72 worker node.

```
$ scp ca.pem 10.10.40.72-key.pem 10.10.40.72.pem 10.10.40.72:~
```

Master nodes

1- Copy the certificates to 10.10.40.60.

```
$ scp ca.pem ca-key.pem kubernetes-key.pem kubernetes.pem 10.10.40.60:~
```

2- Copy the certificates to 10.10.40.61.

```
$ scp ca.pem ca-key.pem kubernetes-key.pem kubernetes.pem 10.10.40.61:~
```



(<http://blog.inkubate.io>)

```
$ scp ca.pem ca-key.pem kubernetes-key.pem kubernetes.pem 10.10.40.62:~
```

Generate the kubeconfig file for the worker nodes

We need to configure the kubelet and the kube-proxy with a kubeconfig file on each worker node to allow them to access the Kubernetes API.

Generate kubeconfig file for the kubelets

1- Add the cluster information for 10.10.40.70.

```
$ kubectl config set-cluster kubernetes \
--certificate-authority=ca.pem \
--embed-certs=true \
--server=https://10.10.40.63:6443 (https://10.10.40.63:6443) \
--kubeconfig=10.10.40.70.kubeconfig
```

2- Add the credentials for 10.10.40.70.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware Sphere
6116 words, 23 min read

```
$ kubectl config set-cluster kubernetes \
--client-certificate=10.10.40.70.pem \
--client-key=10.10.40.70-key.pem \
--embed-certs=true \
--kubeconfig=10.10.40.70.kubeconfig
```

9%

3- Add the context for 10.10.40.70.

```
$ kubectl config set-context default \
--cluster=kubernetes \
--user=system:node:10.10.40.70 \
--kubeconfig=10.10.40.70.kubeconfig
```

4- Use the context for 10.10.40.70.

```
$ kubectl config use-context default --kubeconfig=10.10.40.70.kubeconfig
```

5- Add the cluster information for 10.10.40.71.

```
$ kubectl config set-cluster kubernetes \
--certificate-authority=ca.pem \
--embed-certs=true \
--server=https://10.10.40.63:6443 (https://10.10.40.63:6443) \
--kubeconfig=10.10.40.71.kubeconfig
```

6- Add the credentials for 10.10.40.71.



(<http://blog.inkubate.io>)

```
$ kubectl config set cluster kubernetes \
--client-certificate=10.10.40.71.pem \
--client-key=10.10.40.71-key.pem \
--embed-certs=true \
--kubeconfig=10.10.40.71.kubeconfig
```

7- Add the context for 10.10.40.71.

```
$ kubectl config set-context default \
--cluster=kubernetes \
--user=system:node:10.10.40.71 \
--kubeconfig=10.10.40.71.kubeconfig
```

8- Use the context for 10.10.40.71.

```
$ kubectl config use-context default --kubeconfig=10.10.40.71.kubeconfig
```

9- Add the cluster information for 10.10.40.72.

```
$ kubectl config set-cluster kubernetes \
--certificate-authority=ca.pem \
--embed-certs=true \
--server=https://10.10.40.63:6443 (https://10.10.40.63:6443) \
--kubeconfig=10.10.40.72.kubeconfig
```

10- Add the credentials for 10.10.40.72.

9%



(<http://blog.inkubate.io>)

```
$ kubectl config set-cluster kubernetes \
--client-certificate=10.10.40.72.pem \
--client-key=10.10.40.72-key.pem \
--embed-certs=true \
--kubeconfig=10.10.40.72.kubeconfig
```

11- Add the context for 10.10.40.72.

```
$ kubectl config set-context default \
--cluster=kubernetes \
--user=system:node:10.10.40.72 \
--kubeconfig=10.10.40.72.kubeconfig
```

12- Use the context for 10.10.40.72.

```
$ kubectl config use-context default --kubeconfig=10.10.40.72.kubeconfig
```

Generate the kubeconfig file for the kube-proxies

1- Add the cluster information.

```
$ kubectl config set-cluster kubernetes \
--certificate-authority=ca.pem \
--embed-certs=true \
--server=https://10.10.40.63:6443 (https://10.10.40.63:6443) \
--kubeconfig=kube-proxy.kubeconfig
```

2-Add the credentials.

9%





(http://blog.inkubate.io)

```
$ kubectl config set-cluster default \
--client-certificate=kube-proxy.pem \
--client-key=kube-proxy-key.pem \
--embed-certs=true \
--kubeconfig=kube-proxy.kubeconfig
```

3- Add the context information.

```
$ kubectl config set-context default \
--cluster=kubernetes \
--user=kube-proxy \
--kubeconfig=kube-proxy.kubeconfig
```

4- Use the context.

```
$ kubectl config use-context default --kubeconfig=kube-proxy.kubeconfig
```

Copy the kubeconfig configurations to the worker nodes

1- Copy the configuration to 10.10.40.70.

```
$ scp 10.10.40.70.kubeconfig kube-proxy.kubeconfig 10.10.40.70:~
```

2- Copy the configuration to 10.10.40.71.

```
$ scp 10.10.40.71.kubeconfig kube-proxy.kubeconfig 10.10.40.71:~
```



(<http://blog.inkubate.io>)

```
$ scp 10.10.40.72.kubeconfig kube-proxy.kubeconfig 10.10.40.72:~
```

Generate the data encryption key for secrets

Kubernetes is able to store secrets for us in a key value store. To encrypt this data before storing it, we need an encryption key.

1- Generate a random encryption key.

```
$ ENCRYPTION_KEY=$(head -c 32 /dev/urandom | base64)
```

2- Create a Kubernetes manifest for the encryption configuration.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

kind: EncryptionConfig

apiVersion: v1

resources:

- resources:

- secrets

providers:

- aescbc:

keys:

- name: key1

- secret: \${ENCRYPTION_KEY}

- identity: {}

EOF

3- Copy the manifest to 10.10.40.60.

```
$ scp encryption-config.yaml 10.10.40.60:~
```

4- Copy the manifest to 10.10.40.61.

```
$ scp encryption-config.yaml 10.10.40.61:~
```

5- Copy the manifest to 10.10.40.62.

```
$ scp encryption-config.yaml 10.10.40.62:~
```

9%

One of the main components of the master nodes is the etcd cluster. Etcd is a key value store and is used to store the Kubernetes cluster state. We are going to configure one etcd node per master. With this configuration, the Kubernetes cluster will still be available if one of the etcd fails.

Install etcd on the 10.10.40.60 master node

1- SSH to the 10.10.40.60 virtual machine.

2- Download the etcd binaries.

```
$ wget https://github.com/coreos_(http://github.com/coreos)/etcd/releases/download
```

3- Extract the etcd archive.

```
$ tar xvzf etcd-v3.2.11-linux-amd64.tar.gz
```

4- Move the etcd binaries to /usr/local/bin.

```
$ sudo mv etcd-v3.2.11-linux-amd64/etcd* /usr/local/bin/
```

5- Create the configuration directories.

```
$ sudo mkdir -p /etc/etcd /var/lib/etcd
```



3- Copy the certificates to the /etc/etcd configuration directory.

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

```
$ sudo cp ca.pem kubernetes-key.pem kubernetes.pem /etc/etcd/
```

7- Create an etcd systemd unit file.





(<http://blog.inkubate.io>)

\$ sudo apt-get install etcd

[Unit]

Description=etcd

Documentation=<https://github.com/coreos> (<http://github.com/coreos>).

[Service]

ExecStart=/usr/local/bin/etcd \

--name 10.10.40.60 \

--cert-file=/etc/etcd/kubernetes.pem \

--key-file=/etc/etcd/kubernetes-key.pem \

--peer-cert-file=/etc/etcd/kubernetes.pem \

--peer-key-file=/etc/etcd/kubernetes-key.pem \

--trusted-ca-file=/etc/etcd/ca.pem \

--peer-trusted-ca-file=/etc/etcd/ca.pem \

--peer-client-cert-auth \

--client-cert-auth \

--initial-advertise-peer-urls https://10.10.40.60:2380 \

--listen-peer-urls https://10.10.40.60:2380 \

--listen-client-urls https://10.10.40.60:2379,http://127.0.0.1:2379 \

--advertise-client-urls https://10.10.40.60:2379 \

--initial-cluster-token etcd-cluster-0 \

--initial-cluster 10.10.40.60=https://10.10.40.60:2380,10.10.40.61=h

--initial-cluster-state new \

9%



(<http://blog.inkubate.io>)

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

8- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

9- Enable etcd to start at boot time.

```
$ sudo systemctl enable etcd
```

10- Start etcd.

```
$ sudo systemctl start etcd
```

Install etcd on the 10.10.40.61 master node

1- SSH to the 10.10.40.61 virtual machine.

2- Download the etcd binaries.

```
$ wget https://github.com/coreos/etcd/releases/download/v3.2.11/etcd-v3.2.11-linux-amd64.tar.gz
```

3- Extract the etcd archive.

```
$ tar xvzf etcd-v3.2.11-linux-amd64.tar.gz
```

4- Move the etcd binaries to /usr/local/bin.

```
$ sudo mv etcd-v3.2.11-linux-amd64/etcd* /usr/local/bin/
```

5- Create the configuration directories.

```
$ sudo mkdir -p /etc/etcd /var/lib/etcd
```

6- Copy the certificates to the /etc/etcd configuration directory.



\$ sudo cp kubeconfig-18-from-seattle-on-vmware5.pem /etc/etcd/

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

7- Create an etcd systemd unit file.





(<http://blog.inkubate.io>)

\$ sudo apt-get install etcd -y

[Unit]

Description=etcd

Documentation=<https://github.com/coreos> (<http://github.com/coreos>).

[Service]

ExecStart=/usr/local/bin/etcd \

--name 10.10.40.61 \

--cert-file=/etc/etcd/kubernetes.pem \

--key-file=/etc/etcd/kubernetes-key.pem \

--peer-cert-file=/etc/etcd/kubernetes.pem \

--peer-key-file=/etc/etcd/kubernetes-key.pem \

--trusted-ca-file=/etc/etcd/ca.pem \

--peer-trusted-ca-file=/etc/etcd/ca.pem \

--peer-client-cert-auth \

--client-cert-auth \

--initial-advertise-peer-urls https://10.10.40.61:2380 \

--listen-peer-urls https://10.10.40.61:2380 \

--listen-client-urls https://10.10.40.61:2379,http://127.0.0.1:2379 \

--advertise-client-urls https://10.10.40.61:2379 \

--initial-cluster-token etcd-cluster-0 \

--initial-cluster 10.10.40.60=https://10.10.40.60:2380,10.10.40.61=h

--initial-cluster-state new \

9%



(<http://blog.inkubate.io>)

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

8- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

9- Enable etcd to start at boot time.

```
$ sudo systemctl enable etcd
```

10- Start etcd.

```
$ sudo systemctl start etcd
```

Install etcd on the 10.10.40.62 master node

1- SSH to the 10.10.40.62 virtual machine.

2- Download the etcd binaries.

```
$ wget https://github.com/coreos/(http://github.com/coreos)/etcd/releases/download
```

3- Extract the etcd archive.

```
$ tar xvzf etcd-v3.2.11-linux-amd64.tar.gz
```

4- Move the etcd binaries to /usr/local/bin.

```
$ sudo mv etcd-v3.2.11-linux-amd64/etcd* /usr/local/bin/
```

5- Create the configuration directories.

```
$ sudo mkdir -p /etc/etcd /var/lib/etcd
```

6- Copy the certificates to the /etc/etcd configuration directory.



\$ sudo cp kubeconfig-18-from-seattle-on-vmware5.pem /etc/etcd/

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

7- Create an etcd systemd unit file.





(<http://blog.inkubate.io>)

\$ sudo apt-get install etcd

6116 words - 23 min read.

[Unit]

Description=etcd

Documentation=<https://github.com/coreos> (<http://github.com/coreos>).

[Service]

ExecStart=/usr/local/bin/etcd \

--name 10.10.40.62 \

--cert-file=/etc/etcd/kubernetes.pem \

--key-file=/etc/etcd/kubernetes-key.pem \

--peer-cert-file=/etc/etcd/kubernetes.pem \

--peer-key-file=/etc/etcd/kubernetes-key.pem \

--trusted-ca-file=/etc/etcd/ca.pem \

--peer-trusted-ca-file=/etc/etcd/ca.pem \

--peer-client-cert-auth \

--client-cert-auth \

--initial-advertise-peer-urls https://10.10.40.62:2380 \

--listen-peer-urls https://10.10.40.62:2380 \

--listen-client-urls https://10.10.40.62:2379,http://127.0.0.1:2379 \

--advertise-client-urls https://10.10.40.62:2379 \

--initial-cluster-token etcd-cluster-0 \

--initial-cluster 10.10.40.60=https://10.10.40.60:2380,10.10.40.61=h

--initial-cluster-state new \

9%



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

9%

```
--data-dir=/var/lib/etcd  
Restart=on-failure  
RestartSec=5
```

```
[Install]  
WantedBy=multi-user.target
```

8- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

9- Enable etcd to start at boot time.

```
$ sudo systemctl enable etcd
```

10- Start etcd.

```
$ sudo systemctl start etcd
```

11- Verify that the cluster is up and running.

```
$ ETCDCTL_API=3 etcdctl member list
```

Install the Kubernetes master nodes components

The Kubernetes master nodes need to run three other components: the API, the controller manager, and the scheduler. We will also install a Kubernetes client, kubectl, on each node.

Installation of the 10.10.40.60 master node

1- SSH to 10.10.40.60.

2- Download the binaries.

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```

3- Add the execution permissions to the binaries.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ sudo mv \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl
```

9%

4- Move the binaries to /usr/local/bin.

```
$ sudo mv \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl \
/usr/local/bin/
```

5- Create the API configuration directory.

```
$ sudo mkdir -p /var/lib/kubernetes
```

6- Copy the certificates to the API configuration directory.

```
$ sudo cp \
ca.pem \
ca-key.pem \
kubernetes-key.pem \
kubernetes.pem \
encryption-config.yaml \
/var/lib/kubernetes
```



7- Create an API systemd unit file Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)



 (http://blog.inkubate.io)

\$ sudo systemctl enable kube-apiserver.service

Deploy Kubernetes 1.9 from scratch on VMware vSphere

[Unit]

Description=Kubernetes API Server

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-apiserver \

--admission-control=Initializers,NamespaceLifecycle,NodeRestriction,.

--advertise-address=10.10.40.60 \

--allow-privileged=true \

--apiserver-count=3 \

--audit-log-maxage=30 \

--audit-log-maxbackup=3 \

--audit-log-maxsize=100 \

--audit-log-path=/var/log/audit.log \

--authorization-mode=Node,RBAC \

--bind-address=0.0.0.0 \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--enable-swagger-ui=true \

--etcd-cafile=/var/lib/kubernetes/ca.pem \

--etcd-certfile=/var/lib/kubernetes/kubernetes.pem \

--etcd-keyfile=/var/lib/kubernetes/kubernetes-key.pem \

--etcd-servers=<https://10.10.40.60:2379>,<https://10.10.40.61:2379>,htt

--event-ttl=1h \



(<http://blog.inkubate.io>)

~~--experimental-encryption-provider-config=/var/lib/kubernetes/encryp~~
Deploy Kubernetes 1.9 from scratch on VMware vSphere
6116 words, 23 min read
~~--insecure-bind-address=127.0.0.1 \~~
~~--kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \~~
~~--kubelet-client-certificate=/var/lib/kubernetes/kubernetes.pem \~~
~~--kubelet-client-key=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--kubelet-https=true \~~
~~--runtime-config=api/all \~~
~~--service-account-key-file=/var/lib/kubernetes/ca-key.pem \~~
~~--service-cluster-ip-range=10.32.0.0/24 \~~
~~--service-node-port-range=30000-32767 \~~
~~--tls-ca-file=/var/lib/kubernetes/ca.pem \~~
~~--tls-cert-file=/var/lib/kubernetes/kubernetes.pem \~~
~~--tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--v=2~~

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

9%

8- Create a controller manager systemd unit file.



(http://blog.inkubate.io)

\$ sudo systemctl enable kube-controller-manager.service

Deploy Kubernetes 1.9 from scratch on VMware vSphere

[Unit]

Description=Kubernetes Controller Manager

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-controller-manager \

--address=0.0.0.0 \

--cluster-cidr=10.20.0.0/16 \

--cluster-name=kubernetes \

--cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \

--cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \

--leader-elect=true \

--master=<http://127.0.0.1:8080> (<http://127.0.0.1:8080>) \

--root-ca-file=/var/lib/kubernetes/ca.pem \

--service-account-private-key-file=/var/lib/kubernetes/ca-key.pem \

--service-cluster-ip-range=10.32.0.0/24 \

--v=2

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

9%





(<http://blog.inkubate.io>)

9- Create a scheduler systemd unit file.

```
$ sudo vim /etc/systemd/system/kube-scheduler.service
[Unit]
Description=Kubernetes Scheduler
Documentation=https://github.com/kubernetes/kubernetes (https://github.com/kube

[Service]
ExecStart=/usr/local/bin/kube-scheduler \
  --leader-elect=true \
  --master=http://127.0.0.1:8080 (http://127.0.0.1:8080) \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target
```

10- Reload the daemon configuration.



(<http://blog.inkubate.io>)

6116 words - 23 min read.

11- Enable the services to start at boot time.

```
$ sudo systemctl enable \  
kube-apiserver \  
kube-controller-manager \  
kube-scheduler
```

12- Start the services.

```
$ sudo systemctl start kube-apiserver kube-controller-manager kube-scheduler
```

13- Verify the different components.

```
$ kubectl get componentstatuses
```

Installation of the 10.10.40.61 master node

1- SSH to 10.10.40.61.

2- Download the binaries.

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```



(http://blog.inkubate.io)

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(curl -s https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-amd64.tar.gz -O kube.tar.gz
```

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-amd64.tar.gz -O kube.tar.gz
```

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-amd64.tar.gz -O kube.tar.gz
```

3- Add the execution permissions to the binaries.

```
$ sudo chmod +x \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl
```

4- Move the binaries to /usr/local/bin.

```
$ sudo mv \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl \
/usr/local/bin/
```

5- Create the API configuration directory.

9%





(<http://blog.inkubate.io>)

6- Copy the certificates to the API configuration directory.

```
$ sudo cp \  
ca.pem \  
ca-key.pem \  
kubernetes-key.pem \  
kubernetes.pem \  
encryption-config.yaml \  
/var/lib/kubernetes
```

7- Create an API systemd unit file.



(http://blog.inkubate.io)

\$ sudo systemctl enable kube-apiserver.service

[Unit]

Description=Kubernetes API Server

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-apiserver \

--admission-control=Initializers,NamespaceLifecycle,NodeRestriction,

--advertise-address=10.10.40.61 \

--allow-privileged=true \

--apiserver-count=3 \

--audit-log-maxage=30 \

--audit-log-maxbackup=3 \

--audit-log-maxsize=100 \

--audit-log-path=/var/log/audit.log \

--authorization-mode=Node,RBAC \

--bind-address=0.0.0.0 \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--enable-swagger-ui=true \

--etcd-cafile=/var/lib/kubernetes/ca.pem \

--etcd-certfile=/var/lib/kubernetes/kubernetes.pem \

--etcd-keyfile=/var/lib/kubernetes/kubernetes-key.pem \

--etcd-servers=<https://10.10.40.60:2379>,<https://10.10.40.61:2379>,<https://10.10.40.62:2379> \

--event-ttl=1h \

9%

^



(<http://blog.inkubate.io>)

~~--experimental-encryption-provider-config=/var/lib/kubernetes/encryp~~
Deploy Kubernetes 1.9 from scratch on VMware vSphere
6116 words, 23 min read
~~--insecure-bind-address=127.0.0.1 \~~
~~--kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \~~
~~--kubelet-client-certificate=/var/lib/kubernetes/kubernetes.pem \~~
~~--kubelet-client-key=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--kubelet-https=true \~~
~~--runtime-config=api/all \~~
~~--service-account-key-file=/var/lib/kubernetes/ca-key.pem \~~
~~--service-cluster-ip-range=10.32.0.0/24 \~~
~~--service-node-port-range=30000-32767 \~~
~~--tls-ca-file=/var/lib/kubernetes/ca.pem \~~
~~--tls-cert-file=/var/lib/kubernetes/kubernetes.pem \~~
~~--tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--v=2~~

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

9%

8- Create a controller manager systemd unit file.



(http://blog.inkubate.io)

```
$ sudo systemctl enable kube-controller-manager.service
```

```
[Unit]
```

```
Description=Kubernetes Controller Manager
```

```
Documentation=https://github.com/kubernetes/kubernetes (

```
\[Service\]
```


```

```
ExecStart=/usr/local/bin/kube-controller-manager \
```

```
--address=0.0.0.0 \
```

```
--cluster-cidr=10.20.0.0/16 \
```

```
--cluster-name=kubernetes \
```

```
--cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \
```

```
--cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \
```

```
--leader-elect=true \
```

```
--master=http://127.0.0.1:8080 (http://127.0.0.1:8080) \
```

```
--root-ca-file=/var/lib/kubernetes/ca.pem \
```

```
--service-account-private-key-file=/var/lib/kubernetes/ca-key.pem \
```

```
--service-cluster-ip-range=10.32.0.0/24 \
```

```
--v=2
```

```
Restart=on-failure
```

```
RestartSec=5
```

```
[Install]
```

```
WantedBy=multi-user.target
```

9%





(<http://blog.inkubate.io>)

9- Create a scheduler systemd unit file.

```
$ sudo vim /etc/systemd/system/kube-scheduler.service
[Unit]
Description=Kubernetes Scheduler
Documentation=https://github.com/kubernetes/kubernetes (https://github.com/kube

[Service]
ExecStart=/usr/local/bin/kube-scheduler \
  --leader-elect=true \
  --master=http://127.0.0.1:8080 (http://127.0.0.1:8080) \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target
```

10- Reload the daemon configuration.



(<http://blog.inkubate.io>)

6116 words - 23 min read.

11- Enable the services to start at boot time.

```
$ sudo systemctl enable \  
kube-apiserver \  
kube-controller-manager \  
kube-scheduler
```

12- Start the services.

```
$ sudo systemctl start kube-apiserver kube-controller-manager kube-scheduler
```

13- Verify the different components.

```
$ kubectl get componentstatuses
```

Installation of the 10.10.40.62 master node

1- SSH to 10.10.40.62.

2- Download the binaries.

```
$ wget https://storage.googleapis.com/kubernetes- (https://storage.googleapis.c
```



(http://blog.inkubate.io)

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(curl -s https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-linux-amd64.tar.gz -O kubernetes.tar.gz
```

9%

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-linux-amd64.tar.gz -O kubernetes.tar.gz
```

```
$ wget https://storage.googleapis.com/kubernetes-release/release/$(cat /dev/urandom | tr -dc 'a-z0-9' | fold -w 8 | head -n 1 | tr -d '\n')-linux-amd64.tar.gz -O kubernetes.tar.gz
```

3- Add the execution permissions to the binaries.

```
$ sudo chmod +x \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl
```

4- Move the binaries to /usr/local/bin.

```
$ sudo mv \
kube-apiserver \
kube-controller-manager \
kube-scheduler \
kubectl \
/usr/local/bin/
```

5- Create the API configuration directory.



(<http://blog.inkubate.io>)

6- Copy the certificates to the API configuration directory.

```
$ sudo cp \  
ca.pem \  
ca-key.pem \  
kubernetes-key.pem \  
kubernetes.pem \  
encryption-config.yaml \  
/var/lib/kubernetes
```

7- Create an API systemd unit file.

 (http://blog.inkubate.io)

\$ sudo systemctl enable kube-apiserver.service

Deploy Kubernetes 1.9 from scratch on VMware vSphere

[Unit]

Description=Kubernetes API Server

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-apiserver \

--admission-control=Initializers,NamespaceLifecycle,NodeRestriction,.

--advertise-address=10.10.40.62 \

--allow-privileged=true \

--apiserver-count=3 \

--audit-log-maxage=30 \

--audit-log-maxbackup=3 \

--audit-log-maxsize=100 \

--audit-log-path=/var/log/audit.log \

--authorization-mode=Node,RBAC \

--bind-address=0.0.0.0 \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--enable-swagger-ui=true \

--etcd-cafile=/var/lib/kubernetes/ca.pem \

--etcd-certfile=/var/lib/kubernetes/kubernetes.pem \

--etcd-keyfile=/var/lib/kubernetes/kubernetes-key.pem \

--etcd-servers=<https://10.10.40.60:2379>,<https://10.10.40.61:2379>,<https://10.10.40.62:2379> \

--event-ttl=1h \



(<http://blog.inkubate.io>)

~~--experimental-encryption-provider-config=/var/lib/kubernetes/encryp~~
Deploy Kubernetes 1.9 from scratch on VMware vSphere
6116 words, 23 min read
~~--insecure-bind-address=127.0.0.1 \~~
~~--kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \~~
~~--kubelet-client-certificate=/var/lib/kubernetes/kubernetes.pem \~~
~~--kubelet-client-key=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--kubelet-https=true \~~
~~--runtime-config=api/all \~~
~~--service-account-key-file=/var/lib/kubernetes/ca-key.pem \~~
~~--service-cluster-ip-range=10.32.0.0/24 \~~
~~--service-node-port-range=30000-32767 \~~
~~--tls-ca-file=/var/lib/kubernetes/ca.pem \~~
~~--tls-cert-file=/var/lib/kubernetes/kubernetes.pem \~~
~~--tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \~~
~~--v=2~~

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

9%

8- Create a controller manager systemd unit file.



(http://blog.inkubate.io)

```
$ sudo systemctl enable kube-controller-manager.service
```

```
[Unit]
```

```
Description=Kubernetes Controller Manager
```

```
Documentation=https://github.com/kubernetes/kubernetes (

```
\[Service\]
```


```

```
ExecStart=/usr/local/bin/kube-controller-manager \
```

```
--address=0.0.0.0 \
```

```
--cluster-cidr=10.20.0.0/16 \
```

```
--cluster-name=kubernetes \
```

```
--cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \
```

```
--cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \
```

```
--leader-elect=true \
```

```
--master=http://127.0.0.1:8080 (http://127.0.0.1:8080) \
```

```
--root-ca-file=/var/lib/kubernetes/ca.pem \
```

```
--service-account-private-key-file=/var/lib/kubernetes/ca-key.pem \
```

```
--service-cluster-ip-range=10.32.0.0/24 \
```

```
--v=2
```

```
Restart=on-failure
```

```
RestartSec=5
```

```
[Install]
```

```
WantedBy=multi-user.target
```

9%





(<http://blog.inkubate.io>)

9- Create a scheduler systemd unit file.

```
$ sudo vim /etc/systemd/system/kube-scheduler.service
[Unit]
Description=Kubernetes Scheduler
Documentation=https://github.com/kubernetes/kubernetes (https://github.com/kube

[Service]
ExecStart=/usr/local/bin/kube-scheduler \
  --leader-elect=true \
  --master=http://127.0.0.1:8080 (http://127.0.0.1:8080) \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target
```

10- Reload the daemon configuration.



(<http://blog.inkubate.io>)

11- Enable the services to start at boot time.

```
$ sudo systemctl enable \  
kube-apiserver \  
kube-controller-manager \  
kube-scheduler
```

12- Start the services.

```
$ sudo systemctl start kube-apiserver kube-controller-manager kube-scheduler
```

13- Verify the different components.

```
$ kubectl get componentstatuses
```

Configure the API to access the kubelet

We need to configure each API server to be able to access the kubelet running on each worker node. To achieve this, we need to create a cluster role.

1- SSH to the 10.10.40.60 master node.



2- Create a Kubernetes manifest file kube-apiserver-to-kubelet.yaml.

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

(<http://blog.inkubate.io>)

```
$ vim kube-apiserver-to-kubelet.yaml
apiVersion: rbac.authorization.k8s.io/v1beta1 (http://rbac.authorization.k8s.io)
kind: ClusterRole
metadata:
  annotations:
    rbac.authorization.kubernetes.io/autoupdate (http://rbac.authorization.kubernetes.io/autoupdate): true
  labels:
    kubernetes.io/bootstrapping (http://kubernetes.io/bootstrapping): rbac-defaulting=true
  name: system:kube-apiserver-to-kubelet
rules:
- apiGroups:
  - ""
  resources:
  - nodes/proxy
  - nodes/stats
  - nodes/log
  - nodes/spec
  - nodes/metrics
  verbs:
  - "*"

```

3- Apply the configuration.

```
$ kubectl create -f kube-apiserver-to-kubelet.yaml
```



4- Create a manifest file to bind the cluster role to the kubernetes user used by the API server.

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

(<http://blog.inkubate.io>)

9%

```
$ vim kube-apiserver-to-kubelet-bind.yaml
apiVersion: rbac.authorization.k8s.io/v1beta1 (http://rbac.authorization.k8s.io)
kind: ClusterRoleBinding
metadata:
  name: system:kube-apiserver
  namespace: ""
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:kube-apiserver-to-kubelet
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: User
  name: kubernetes
```

5- Apply the configuration.

```
$ kubectl create -f kube-apiserver-to-kubelet-bind.yaml
```

Verify that you can access the API server through the HAProxy load balancer

At this point, you should be able to access the API server through the HAProxy load balancer.



(<http://blog.inkubate.io>)

\$ curl -s https://raw.githubusercontent.com/inkubate/inkubate/master/6443/version

6116 words - 23 min read.

9%

Install the Kubernetes worker nodes components

The Kubernetes worker nodes need to run various components. We will install the cni plugin, cri containerd, the kube-proxy, the kubelet, and the kubectl client on each nodes.

Installation of the 10.10.40.70 worker node

1- SSH to 10.10.40.70.

2- Disable the swap as the kubelet refuses to start with the swap enabled.

```
$ sudo swapoff -a
```

```
$ sudo sed -i '/ swap / s/^/#/' /etc/fstab
```

3- Modify the hostname to match the IP.

```
$ sudo hostnamectl set-hostname 10.10.40.70
```

4- Update the virtual machine.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

```
$ sudo apt-get upgrade
```

5- Install socat. Socat is a dependency of the kubectl port-forward command. The port forwarding will not work if you don't install socat.

```
$ sudo apt-get install socat
```

6- Download the binaries.

```
$ wget https://github.com/container networking_(http://github.com/container netwo
```

```
$ wget https://github.com/kubernetes_(http://github.com/kubernetes)-incubator/c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

7- Create the installation directories.





(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

```
$ sudo cp /etc/cni/net.d \
/opt/cni/bin \
/var/lib/kubelet \
/var/lib/kube-proxy \
/var/lib/kubernetes \
/var/run/kubernetes
```

8- Install the binaries in their respective directory.

```
$ sudo tar -xvzf cni-plugins-amd64-v0.6.0.tgz -C /opt/cni/bin/
```

```
$ sudo tar -xvzf cri-containerd-1.0.0-beta.0.linux-amd64.tar.gz -C /
```

```
$ chmod +x kubect1 kube-proxy kubelet
```

```
$ sudo mv kubect1 kube-proxy kubelet /usr/local/bin/
```

9- Configure the CNI networking.

9%



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ sudo vim /etc/cni/net.d/10-bridge.conf
{
  "cniVersion": "0.3.1",
  "name": "bridge",
  "type": "bridge",
  "bridge": "cnio0",
  "isGateway": true,
  "ipMasq": true,
  "ipam": {
    "type": "host-local",
    "ranges": [
      [{"subnet": "10.20.0.0/24"}]
    ],
    "routes": [{"dst": "0.0.0.0/0"}]
  }
}
```

```
$ sudo vim /etc/cni/net.d/99-loopback.conf
{
  "cniVersion": "0.3.1",
  "type": "loopback"
}
```

10- Configure the kubelet.

```
$ sudo cp 10.10.40.70-key.pem 10.10.40.70.pem /var/lib/kubelet
```

```
$ sudo cp 10.10.40.70.kubeconfig /var/lib/kubelet/kubeconfig
```

9%



\$ sudo curl -s https://raw.githubusercontent.com/inkubate-io/inkubate/master/vsphere/

6116 words - 23 min read.

9%

(http://blog.inkubate.io)

11- Create the kubelet systemd unit file.





(http://blog.inkubate.io)

\$ sudo systemctl enable kubelet.service

[Unit]

Description=Kubernetes Kubelet

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

After=cri-containerd.service

Requires=cri-containerd.service

[Service]

ExecStart=/usr/local/bin/kubelet \

--allow-privileged=true \

--anonymous-auth=false \

--authorization-mode=Webhook \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--cloud-provider= \

--cluster-dns=10.32.0.10 \

--cluster-domain=cluster.local \

--container-runtime=remote \

--container-runtime-endpoint=unix:///var/run/cri-containerd.sock \

--image-pull-progress-deadline=2m \

--kubeconfig=/var/lib/kubelet/kubeconfig \

--network-plugin=cni \

--pod-cidr=10.20.0.0/24 \

--register-node=true \

--runtime-request-timeout=15m \

9%





(<http://blog.inkubate.io>)

```
--tls-cert-file=/var/lib/kubelet/10.10.40.70.pem \  
--tls-private-key-file=/var/lib/kubelet/10.10.40.70-key.pem \  
--v=2  
Restart=on-failure  
RestartSec=5
```

[Install]

WantedBy=multi-**user**.target

12- Configure the kube-proxy.

```
$ sudo cp kube-proxy.kubeconfig /var/lib/kube-proxy/kubeconfig
```

13- Create the kube-proxy systemd unit file.

9%



(http://blog.inkubate.io)

\$ sudo systemctl enable kube-proxy.service

[Unit]

Description=Kubernetes Kube Proxy

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-proxy \

--cluster-cidr=10.20.0.0/16 \

--kubeconfig=/var/lib/kube-proxy/kubeconfig \

--proxy-mode=iptables \

--v=2

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

14- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

15- Enable the services to start at boot time.

9%





(<http://blog.inkubate.io>)

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

6116 words - 23 min read.

9%

16- Start the services.

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

Installation of the 10.10.40.71 worker node

1- SSH to 10.10.40.71.

2- Disable the swap as the kubelet refuses to start with the swap enabled.

```
$ sudo swapoff -a
```

```
$ sudo sed -i '/ swap / s/^/#/' /etc/fstab
```

3- Modify the hostname to match the IP.

```
$ sudo hostnamectl set-hostname 10.10.40.71
```

4- Update the virtual machine.

```
$ sudo apt-get update
```

```
$ sudo apt-get upgrade
```



5- Install socat. Socat is a dependency of the kubectl port-forward command. The port forwarding will not work if you don't install socat.

```
$ sudo apt-get install socat
```

6- Download the binaries.

```
$ wget https://github.com/container networking_(http://github.com/container netwo
```

```
$ wget https://github.com/kubernetes_(http://github.com/kubernetes)-incubator/c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

7- Create the installation directories.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ sudo cp -r \
/etc/cni/net.d \
/opt/cni/bin \
/var/lib/kubelet \
/var/lib/kube-proxy \
/var/lib/kubernetes \
/var/run/kubernetes
```

9%

8- Install the binaries in their respective directory.

```
$ sudo tar -xvzf cni-plugins-amd64-v0.6.0.tgz -C /opt/cni/bin/
```

```
$ sudo tar -xvzf cri-containerd-1.0.0-beta.0.linux-amd64.tar.gz -C /
```

```
$ chmod +x kubect1 kube-proxy kubelet
```

```
$ sudo mv kubect1 kube-proxy kubelet /usr/local/bin/
```

9- Configure the CNI networking.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ sudo vim /etc/cni/net.d/10-bridge.conf
{
  "cniVersion": "0.3.1",
  "name": "bridge",
  "type": "bridge",
  "bridge": "cnio0",
  "isGateway": true,
  "ipMasq": true,
  "ipam": {
    "type": "host-local",
    "ranges": [
      [{"subnet": "10.20.1.0/24"}]
    ],
    "routes": [{"dst": "0.0.0.0/0"}]
  }
}
```

```
$ sudo vim /etc/cni/net.d/99-loopback.conf
{
  "cniVersion": "0.3.1",
  "type": "loopback"
}
```

10- Configure the kubelet.

```
$ sudo cp 10.10.40.71-key.pem 10.10.40.71.pem /var/lib/kubelet
```

```
$ sudo cp 10.10.40.71.kubeconfig /var/lib/kubelet/kubeconfig
```

9%



\$ sudo curl -s https://raw.githubusercontent.com/inkubate-io/inkubate/master/vsphere/

6116 words - 23 min read.

9%

(http://blog.inkubate.io)

11- Create the kubelet systemd unit file.





(http://blog.inkubate.io)

\$ sudo systemctl enable kubelet.service

[Unit]

Description=Kubernetes Kubelet

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

After=cri-containerd.service

Requires=cri-containerd.service

[Service]

ExecStart=/usr/local/bin/kubelet \

--allow-privileged=true \

--anonymous-auth=false \

--authorization-mode=Webhook \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--cloud-provider= \

--cluster-dns=10.32.0.10 \

--cluster-domain=cluster.local \

--container-runtime=remote \

--container-runtime-endpoint=unix:///var/run/cri-containerd.sock \

--image-pull-progress-deadline=2m \

--kubeconfig=/var/lib/kubelet/kubeconfig \

--network-plugin=cni \

--pod-cidr=10.20.1.0/24 \

--register-node=true \

--runtime-request-timeout=15m \

9%





(<http://blog.inkubate.io>)

```
--tls-cert-file=/var/lib/kubelet/10.10.40.71.pem \  
--tls-private-key-file=/var/lib/kubelet/10.10.40.71-key.pem \  
--v=2  
Restart=on-failure  
RestartSec=5
```

[Install]

WantedBy=multi-**user**.target

12- Configure the kube-proxy.

```
$ sudo cp kube-proxy.kubeconfig /var/lib/kube-proxy/kubeconfig
```

13- Create the kube-proxy systemd unit file.

9%



(http://blog.inkubate.io)

\$ sudo systemctl enable kube-proxy.service

[Unit]

Description=Kubernetes Kube Proxy

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-proxy \

--cluster-cidr=10.20.0.0/16 \

--kubeconfig=/var/lib/kube-proxy/kubeconfig \

--proxy-mode=iptables \

--v=2

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

14- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

15- Enable the services to start at boot time.

9%





(<http://blog.inkubate.io>)

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

6116 words - 23 min read.

9%

16- Start the services.

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

Installation of the 10.10.40.72 worker node

1- SSH to 10.10.40.72.

2- Disable the swap as the kubelet refuses to start with the swap enabled.

```
$ sudo swapoff -a
```

```
$ sudo sed -i '/ swap / s/^/#/' /etc/fstab
```

3- Modify the hostname to match the IP.

```
$ sudo hostnamectl set-hostname 10.10.40.72
```

4- Update the virtual machine.

```
$ sudo apt-get update
```

```
$ sudo apt-get upgrade
```



5- Install socat. Socat is a dependency of the kubectl port-forward command. The port forwarding will not work if you don't install socat.

(<http://blog.inkubate.io>)

```
$ sudo apt-get install socat
```

6- Download the binaries.

```
$ wget https://github.com/containernetworking_(http://github.com/containernetwo
```

```
$ wget https://github.com/kubernetes_(http://github.com/kubernetes)-incubator/c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

```
$ wget https://storage.googleapis.com/kubernetes-_(https://storage.googleapis.c
```

7- Create the installation directories.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

```
$ sudo cp /etc/cni/net.d \
/opt/cni/bin \
/var/lib/kubelet \
/var/lib/kube-proxy \
/var/lib/kubernetes \
/var/run/kubernetes
```

8- Install the binaries in their respective directory.

```
$ sudo tar -xvzf cni-plugins-amd64-v0.6.0.tgz -C /opt/cni/bin/
```

```
$ sudo tar -xvzf cri-containerd-1.0.0-beta.0.linux-amd64.tar.gz -C /
```

```
$ chmod +x kubect1 kube-proxy kubelet
```

```
$ sudo mv kubect1 kube-proxy kubelet /usr/local/bin/
```

9- Configure the CNI networking.

9%



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ sudo vim /etc/cni/net.d/10-bridge.conf
{
  "cniVersion": "0.3.1",
  "name": "bridge",
  "type": "bridge",
  "bridge": "cnio0",
  "isGateway": true,
  "ipMasq": true,
  "ipam": {
    "type": "host-local",
    "ranges": [
      [{"subnet": "10.20.2.0/24"}]
    ],
    "routes": [{"dst": "0.0.0.0/0"}]
  }
}
```

```
$ sudo vim /etc/cni/net.d/99-loopback.conf
{
  "cniVersion": "0.3.1",
  "type": "loopback"
}
```

10- Configure the kubelet.

```
$ sudo cp 10.10.40.72-key.pem 10.10.40.72.pem /var/lib/kubelet
```

```
$ sudo cp 10.10.40.72.kubeconfig /var/lib/kubelet/kubeconfig
```

9%



\$ sudo apt-get install kubelet kubeadm kubectl

6116 words - 23 min read.

9%

(http://blog.inkubate.io)

11- Create the kubelet systemd unit file.





(http://blog.inkubate.io)

\$ sudo systemctl enable kubelet.service

[Unit]

Description=Kubernetes Kubelet

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

After=cri-containerd.service

Requires=cri-containerd.service

[Service]

ExecStart=/usr/local/bin/kubelet \

--allow-privileged=true \

--anonymous-auth=false \

--authorization-mode=Webhook \

--client-ca-file=/var/lib/kubernetes/ca.pem \

--cloud-provider= \

--cluster-dns=10.32.0.10 \

--cluster-domain=cluster.local \

--container-runtime=remote \

--container-runtime-endpoint=unix:///var/run/cri-containerd.sock \

--image-pull-progress-deadline=2m \

--kubeconfig=/var/lib/kubelet/kubeconfig \

--network-plugin=cni \

--pod-cidr=10.20.2.0/24 \

--register-node=true \

--runtime-request-timeout=15m \

9%





(<http://blog.inkubate.io>)

```
--tls-cert-file=/var/lib/kubelet/10.10.40.72.pem \  
--tls-private-key-file=/var/lib/kubelet/10.10.40.72-key.pem \  
--v=2  
Restart=on-failure  
RestartSec=5
```

[Install]

WantedBy=multi-**user**.target

12- Configure the kube-proxy.

```
$ sudo cp kube-proxy.kubeconfig /var/lib/kube-proxy/kubeconfig
```

13- Create the kube-proxy systemd unit file.

9%



(http://blog.inkubate.io)

\$ sudo systemctl enable kube-proxy.service

[Unit]

Description=Kubernetes Kube Proxy

Documentation=<https://github.com/kubernetes/kubernetes> (<https://github.com/kube>

[Service]

ExecStart=/usr/local/bin/kube-proxy \

--cluster-cidr=10.20.0.0/16 \

--kubeconfig=/var/lib/kube-proxy/kubeconfig \

--proxy-mode=iptables \

--v=2

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target

14- Reload the daemon configuration.

```
$ sudo systemctl daemon-reload
```

15- Enable the services to start at boot time.

9%





(<http://blog.inkubate.io>)

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

6116 words - 23 min read.

9%

16- Start the services.

```
$ sudo systemctl start containerd cri-containerd kubelet kube-proxy
```

Verify that the Kubernetes cluster is up and running

1- SSH to 10.10.40.60.

2- List the worker nodes.

```
$ kubectl get nodes
```

The three worker nodes should be listed and in a ready state.

Configure your client machine to access the Kubernetes cluster



We need to create a kubeconfig file on the client machine to be able to manage the Kubernetes cluster with kubectl from this machine.

(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words, 22 min read

9%

1- Add the cluster information.

```
$ kubectl config set-cluster kubernetes \  
--certificate-authority=ca.pem \  
--embed-certs=true \  
--server=https://10.10.40.63:6443 (https://10.10.40.63:6443)
```

2- Add the credentials.

```
$ kubectl config set-credentials admin \  
--client-certificate=admin.pem \  
--client-key=admin-key.pem
```

3- Add the context.

```
$ kubectl config set-context kubernetes --cluster=kubernetes --user=admin
```

4- Use the context.

```
$ kubectl config use-context kubernetes
```

5- Verify that you can access the Kubernetes cluster from the client machine.

```
$ kubectl get nodes
```



Configuring the networking route between pods

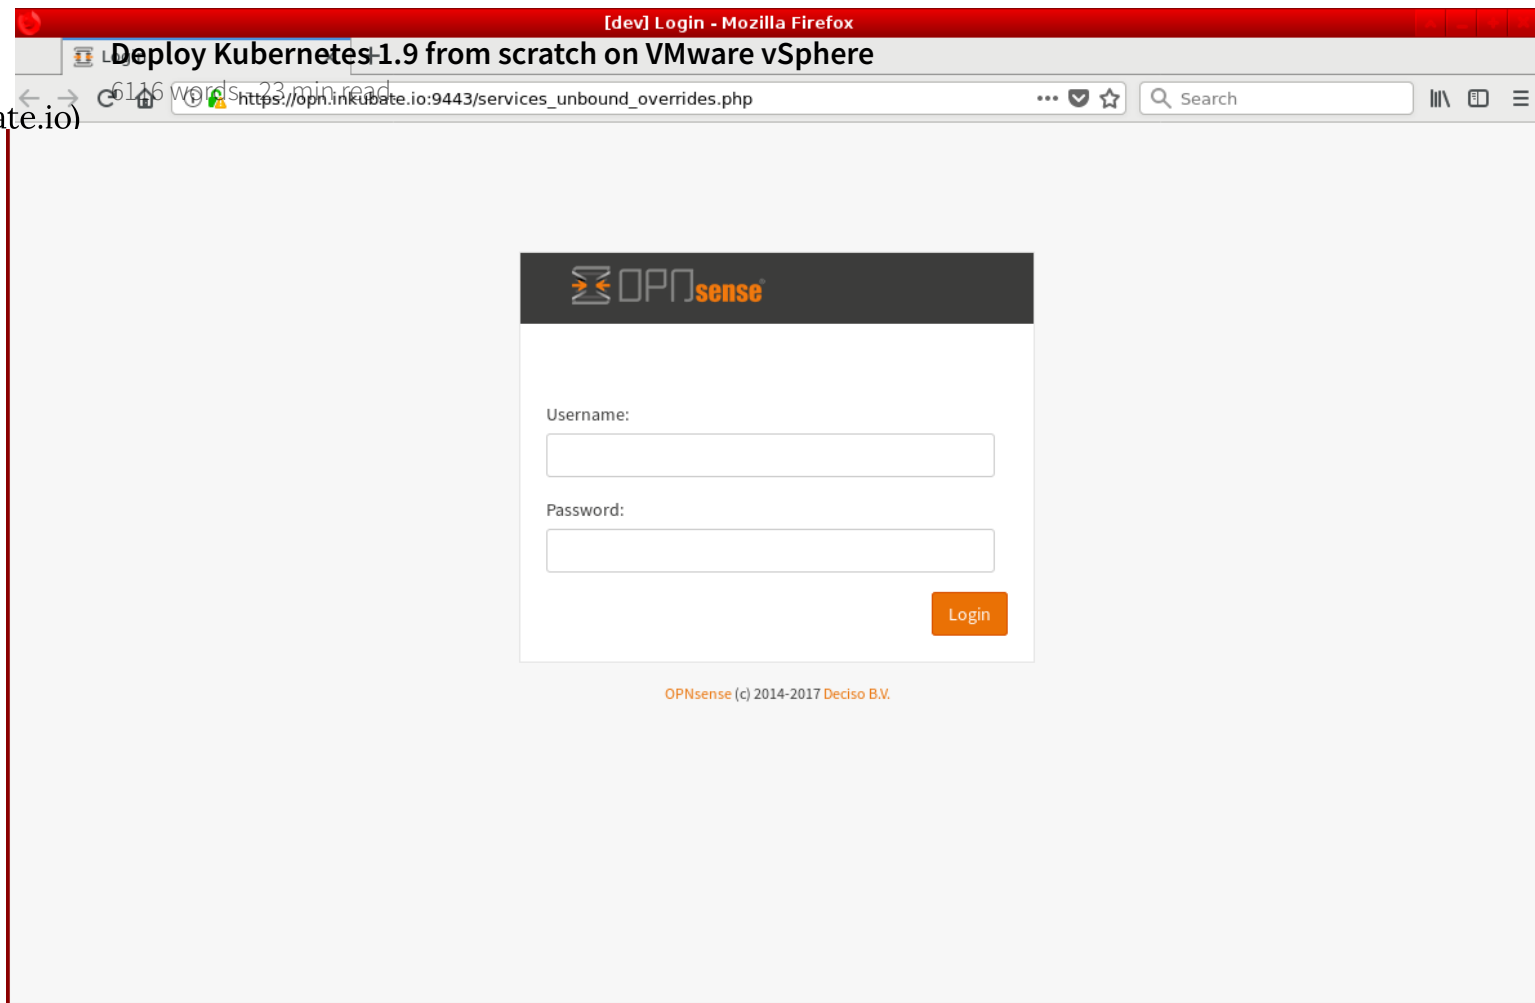
The kubernetes cluster is now up and running, but we still have one issue to solve. Each worker node uses a different network to assign an IP to the pods. It means that two pods located on two different worker nodes cannot communicate between them. To solve this issue, we need to configure static routes on the gateway of the network 10.10.40.0/24. In my case, the gateway is 10.10.40.1 and is an OPNSense server.

If you don't have access to your network gateway, you can still solve this problem by creating an overlay network for the pods with something like flannel, weave net, or calico.

1- Login to your OPNSense gateway.



(<http://blog.inkubate.io>)



9%

2- Add 10.10.40.70 as a new gateway.



6116 words - 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

All

Status

Group

Group Status

Log File

Routes

High Availability

Configuration

Trust

Wizard

Log File







Diagnostics

Interfaces

System: Gateways: All

Add gateway

The changes have been applied successfully.

	Name	Interface	Gateway	Monitor IP	Description	
<input checked="" type="checkbox"/>	WAN_DHCP (default)	WAN	213.239.214.65	213.239.214.65	Interface WAN_DHCP Gateway	 
<input type="checkbox"/>	NSX_Gateway	LAB_STATIC	10.10.40.150			  
						

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_gateways_edit.php

9%

^

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_gateways_edit.php

OPNsense root@opn.inkubate.io

Lobby
Reporting
System
Firmware
Access
Settings
Gateways
All
Status
Group
Group Status
Log File
Routes
High Availability
Configuration
Trust
Wizard
Log File
Diagnostics
Interfaces
Firewall

Address Family: IPv4
Name: K8S_WORKER_0
Gateway: 10.10.40.70
Default Gateway: ☐
Far Gateway: ☐
Disable Gateway Monitoring: ☒
Monitor IP:
Mark Gateway as Down: ☐
Advanced: Advanced - Show advanced option
Description:
Save Cancel

OPNsense (c) 2014-2017 Deciso B.V.

9%

3- Add 10.10.40.71 as a new gateway.

^

6116 words - 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

All

Status

Group

Group Status

Log File

Routes

High Availability

Configuration

Trust

Wizard

Log File










Diagnostics

Interfaces

System: Gateways: All

Add gateway

The changes have been applied successfully.

	Name	Interface	Gateway	Monitor IP	Description	
<input checked="" type="checkbox"/>	WAN_DHCP (default)	WAN	213.239.214.65	213.239.214.65	Interface WAN_DHCP Gateway	 
<input type="checkbox"/>	NSX_Gateway	LAB_STATIC	10.10.40.150			  
<input type="checkbox"/>	K8S_WORKER_0	LAB_STATIC	10.10.40.70			  
						

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_gateways_edit.php

9%

^

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_gateways_edit.php

OPNsense root@opn.inkubate.io

Lobby
Reporting
System
Firmware
Access
Settings
Gateways
All
Status
Group
Group Status
Log File
Routes
High Availability
Configuration
Trust
Wizard
Log File
Diagnostics
Interfaces
Firewall

Address Family: IPv4

Name: K8S_WORKER_1

Gateway: 10.10.40.71

Default Gateway: ☐

Far Gateway: ☐

Disable Gateway Monitoring: ☒

Monitor IP:

Mark Gateway as Down: ☐

Advanced: Advanced - Show advanced option

Description:

Save Cancel

OPNsense (c) 2014-2017 Deciso B.V.

9%

4- Add 10.10.40.72 as a new gateway.

6116 words, 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words, 23 min read

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

All

Status

Group

Group Status

Log File

Routes

High Availability

Configuration

Trust

Wizard

Log File

Diagnostics













Interfaces

System: Gateways: All

Add gateway

The gateway configuration has been changed.
You must apply the changes in order for them to take effect.

Apply changes

	Name	Interface	Gateway	Monitor IP	Description	
	WAN_DHCP (default)	WAN	213.239.214.65	213.239.214.65	Interface WAN_DHCP Gateway	 
<input type="checkbox"/>	NSX_Gateway	LAB_STATIC	10.10.40.150			  
<input type="checkbox"/>	K8S_WORKER_0	LAB_STATIC	10.10.40.70			  
<input type="checkbox"/>	K8S_WORKER_1	LAB_STATIC	10.10.40.71			  
						

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_gateways_edit.php

9%

^

6116 words - 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_gateways_edit.php

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

All

Status

Group

Group Status

Log File

Routes

High Availability

Configuration

Trust

Wizard

Log File

Diagnostics

Interfaces

Firewall

Address Family

IPv4

Name

K8S_WORKER_2

Gateway

10.10.40.72

Default Gateway

☐

Far Gateway

☐

Disable Gateway Monitoring

☒

Monitor IP

Mark Gateway as Down

☐

Advanced

Advanced - Show advanced option

Description

Save Cancel

OPNsense (c) 2014-2017 Deciso B.V.

5- Apply the changes.

9%

6116 words - 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

All

Status

Group

Group Status

Log File

Routes

High Availability

Configuration

Trust

Wizard

Log File

Diagnostics

Interfaces

















Firewall

System: Gateways: All

Add gateway

The gateway configuration has been changed.
You must apply the changes in order for them to take effect.

Apply changes

	Name	Interface	Gateway	Monitor IP	Description	
<input checked="" type="checkbox"/>	WAN_DHCP (default)	WAN	213.239.214.65	213.239.214.65	Interface WAN_DHCP Gateway	  
<input type="checkbox"/>	NSX_Gateway	LAB_STATIC	10.10.40.150			  
<input type="checkbox"/>	K8S_WORKER_0	LAB_STATIC	10.10.40.70			  
<input type="checkbox"/>	K8S_WORKER_1	LAB_STATIC	10.10.40.71			  
<input type="checkbox"/>	K8S_WORKER_2	LAB_STATIC	10.10.40.72			  
						

OPNsense (c) 2014-2017 Deciso B.V.

6- Create a new route for the pods running on 10.10.40.70.

9%

6116 words - 23 min read
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

OPNsense

root@opn.inkubate.io

Lobby

Reporting

System

- Firmware
- Access
- Settings
- Gateways
- Routes
 - All
 - Status
 - Log File
- High Availability
- Configuration
- Trust
- Wizard
- Log File
- Diagnostics

Interfaces

Firewall

VPN

System: Routes: All

Add route

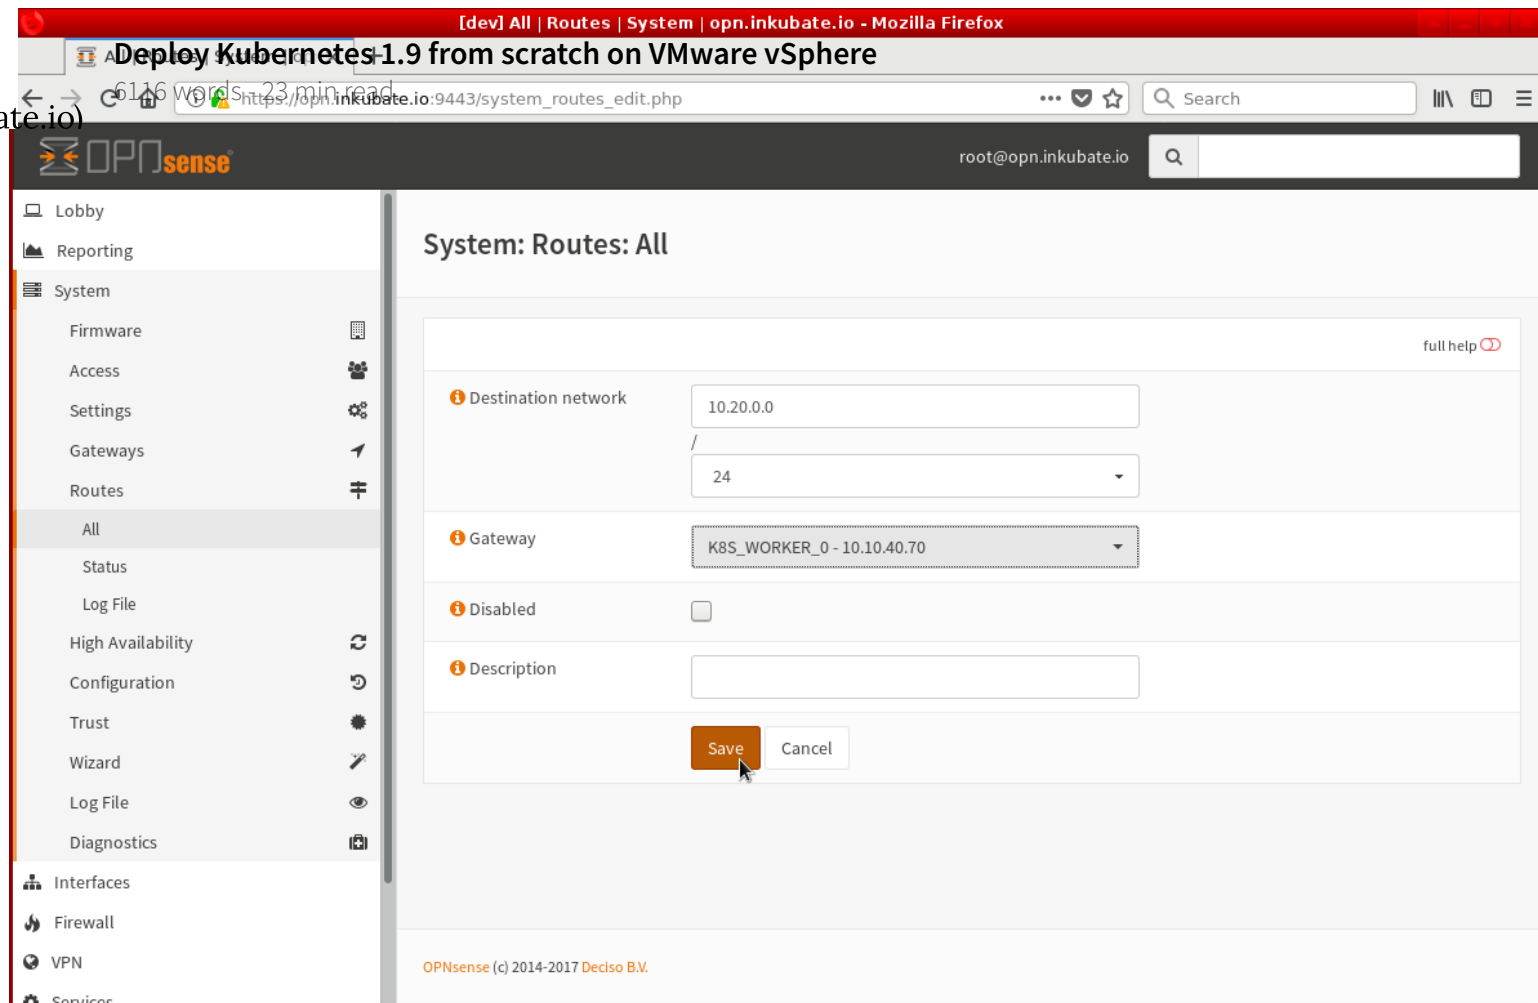
	Network	Gateway	Interface	Description
<input type="checkbox"/>	192.168.0.0/16	NSX_Gateway - 10.10.40.150	LAB_STATIC	

Do not enter static routes for networks assigned on any interface of this firewall. Static routes are only used for networks reachable via a different router, and not reachable via your default gateway.

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_routes_edit.php

←
(http://blog.inkubate.io)



7- Create a new route for the pods running on 10.10.40.71.

9%

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_routes.php

Search

OPNsense

root@opn.inkubate.io

Q

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

Routes

All

Status

Log File

High Availability

Configuration

Trust

Wizard

Log File

Diagnostics

Interfaces

Firewall

VPN

System: Routes: All

Add route

The static route configuration has been changed.
You must apply the changes in order for them to take effect.

Apply changes

	Network	Gateway	Interface	Description	
<input type="checkbox"/>	▶ 192.168.0.0/16	NSX_Gateway - 10.10.40.150	LAB_STATIC		<div>← ↗ 🗑️ 📄</div>
<input type="checkbox"/>	▶ 10.20.0.0/24	K8S_WORKER_0 - 10.10.40.70	LAB_STATIC		<div>← ↗ 🗑️ 📄</div>
					<div>← 🗑️ +</div>

Do not enter static routes for networks assigned on any interface of this firewall. Static routes are only used for networks reachable via a different router, and not reachable via your default gateway.

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_routes_edit.php

9%

^

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_routes_edit.php

OPNsense

root@opn.inkubate.io

System: Routes: All

full help ⓘ

Destination network 10.20.1.0 / 24

Gateway K8S_WORKER_1 - 10.10.40.71

Disabled ☐

Description

Save Cancel

OPNsense (c) 2014-2017 Deciso B.V.

8- Create a new route for the pods running on 10.10.40.72.

9%

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_routes.php

Search

OPNsense

root@opn.inkubate.io

Q

Lobby

Reporting

System

Firmware

Access

Settings

Gateways

Routes

All

Status

Log File

High Availability

Configuration

Trust

Wizard

Log File

Diagnostics

Interfaces

Firewall

VPN

System: Routes: All

Add route

The static route configuration has been changed.
You must apply the changes in order for them to take effect.

Apply changes

	Network	Gateway	Interface	Description	
<input type="checkbox"/>	▶ 192.168.0.0/16	NSX_Gateway - 10.10.40.150	LAB_STATIC		← ✎ 🗑 📄
<input type="checkbox"/>	▶ 10.20.0.0/24	K8S_WORKER_0 - 10.10.40.70	LAB_STATIC		← ✎ 🗑 📄
<input type="checkbox"/>	▶ 10.20.1.0/24	K8S_WORKER_1 - 10.10.40.71	LAB_STATIC		← ✎ 🗑 📄
					← 🗑 +

Do not enter static routes for networks assigned on any interface of this firewall. Static routes are only used for networks reachable via a different router, and not reachable via your default gateway.

OPNsense (c) 2014-2017 Deciso B.V.

https://opn.inkubate.io:9443/system_routes_edit.php

9%

^

←
(http://blog.inkubate.io)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

https://opn.inkubate.io:9443/system_routes_edit.php

OPNsense

root@opn.inkubate.io

System: Routes: All

full help

Destination network

10.20.2.0

/

24

Gateway

K8S_WORKER_2 - 10.10.40.72

Disabled

☐

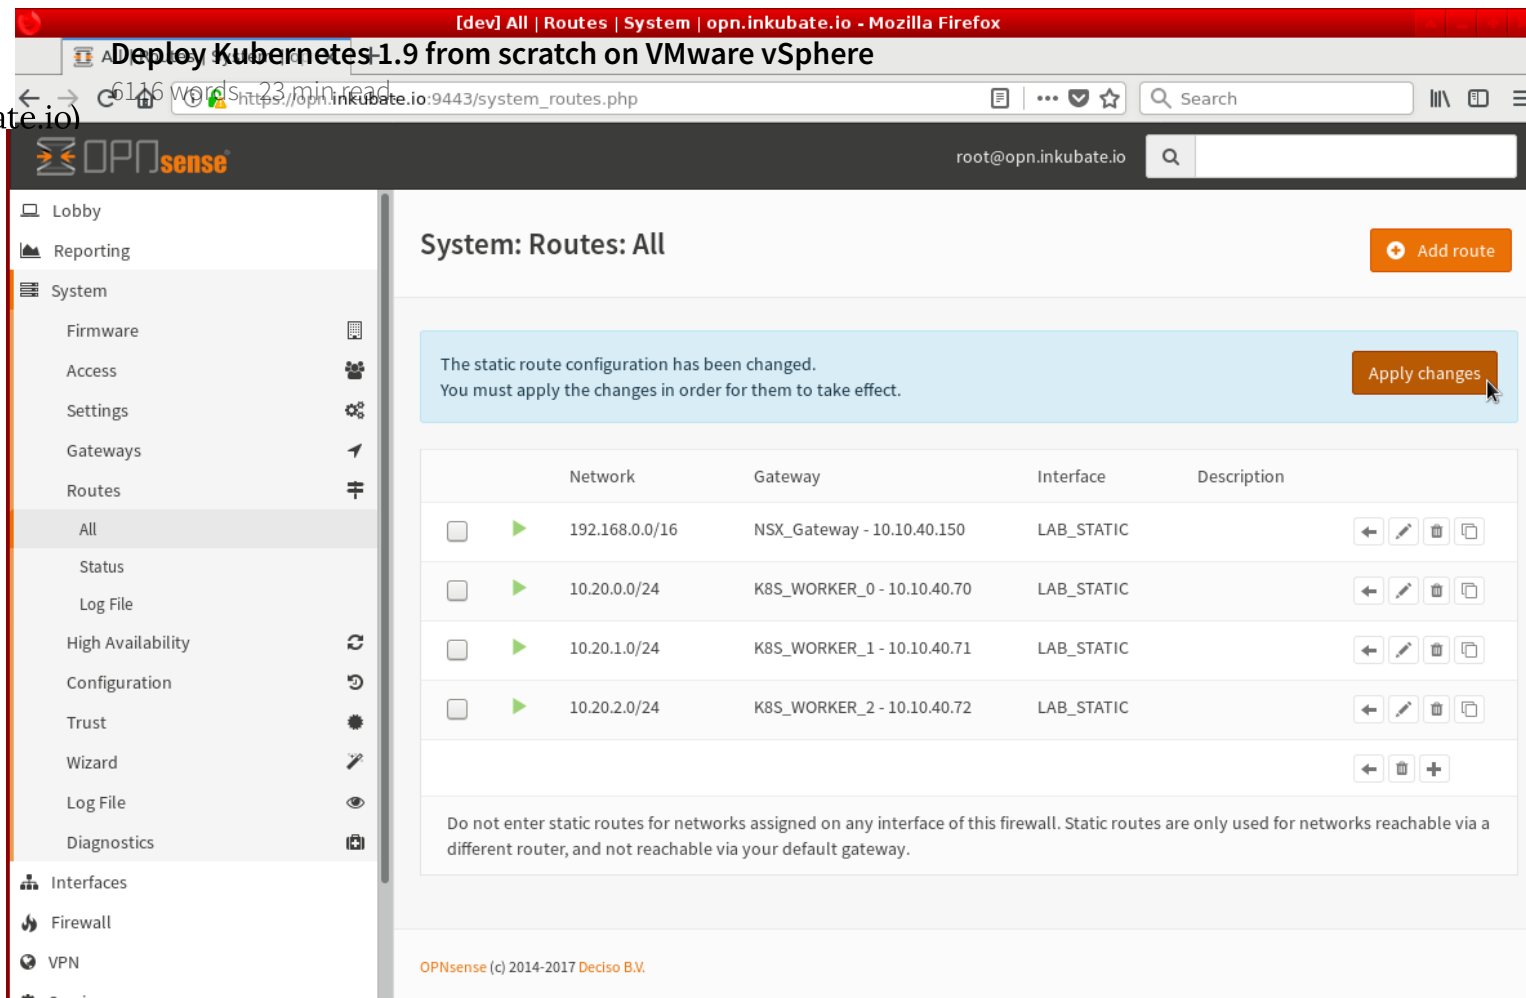
Description

Save Cancel

OPNsense (c) 2014-2017 Deciso B.V.

9- Apply the changes.

9%



Installing Kubernetes add-ons

We will deploy three Kubernetes add-ons on our new cluster. The kube-dns add-on which provides a DNS based service discovery for the pods running in the cluster, the dashboard add-on to have a graphical view of the cluster and the Heapster add-on to monitor our workload.



Installing kube-dns from scratch on VMware vSphere

6116 words - 23 min read.

(<http://blog.inkubate.io>)

I- Create the kube-dns manifest.

9%





(<http://blog.inkubate.io>)

\$ vi DeployKubernetes1.9 from scratch on VMware vSphere

apiVersion: v1

kind: Service

metadata:

name: kube-dns

namespace: kube-system

labels:

k8s-app: kube-dns

kubernetes.io/cluster- (http://kubernetes.io/cluster-).service: "true"

addonmanager.kubernetes.io/mode: Reconcile

kubernetes.io/name: "KubeDNS"

spec:

selector:

k8s-app: kube-dns

clusterIP: 10.32.0.10

ports:

- name: dns

port: 53

protocol: UDP

- name: dns-tcp

port: 53

protocol: TCP

apiVersion: v1

kind: ServiceAccount

metadata:

name: kube-dns

namespace: kube-system

labels:

kubernetes.io/cluster- (http://kubernetes.io/cluster-).service: "true"

addonmanager.kubernetes.io/mode: Reconcile

9%





(<http://blog.inkubate.io>)

--- Deploy Kubernetes 1.9 from scratch on VMware vSphere

apiVersion: v1

0116 words - 23 min read.

kind: ConfigMap

metadata:

name: kube-dns

namespace: kube-system

labels:

addonmanager.kubernetes.io/mode: EnsureExists

apiVersion: extensions/v1beta1

kind: Deployment

metadata:

name: kube-dns

namespace: kube-system

labels:

k8s-app: kube-dns

kubernetes.io/cluster- (<http://kubernetes.io/cluster->).service: "true"

addonmanager.kubernetes.io/mode: Reconcile

spec:

replicas: not specified here:

1. In order to make Addon Manager do not reconcile this replicas parameter.

2. Default is 1.

3. Will be tuned in real time if DNS horizontal auto-scaling is turned on.

strategy:

rollingUpdate:

maxSurge: 10%

maxUnavailable: 0

selector:

matchLabels:

k8s-app: kube-dns

template:

9%





(<http://blog.inkubate.io>)

metadata: Deploy Kubernetes 1.9 from scratch on VMware vSphere

labels:
6116 words - 23 min read.

9%

```
  k8s-app: kube-dns
  annotations:
    scheduler.alpha.kubernetes.io/critical- (http://scheduler.alpha.kuberne
spec:
  tolerations:
    - key: "CriticalAddonsOnly"
      operator: "Exists"
  volumes:
    - name: kube-dns-config
      configMap:
        name: kube-dns
        optional: true
  containers:
    - name: kubedns
      image: k8s.gcr.io/k8s-dns-kube-dns- (http://k8s.gcr.io/k8s-dns-kube-dns
      resources:
        # TODO: Set memory limits when we've profiled the container for large
        # clusters, then set request = limit to keep this container in
        # guaranteed class. Currently, this container falls into the
        # "burstable" category so the kubelet doesn't backoff from restarting
        limits:
          memory: 170Mi
        requests:
          cpu: 100m
          memory: 70Mi
  livenessProbe:
    httpGet:
      path: /healthcheck/kubedns
      port: 10054
```



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
    scheme: HTTP
    initialDelaySeconds: 60
    timeoutSeconds: 5
    successThreshold: 1
    failureThreshold: 5
  readinessProbe:
    httpGet:
      path: /readiness
      port: 8081
      scheme: HTTP
      # we poll on pod startup for the Kubernetes master service and
      # only setup the /readiness HTTP server once that's available.
    initialDelaySeconds: 3
    timeoutSeconds: 5
  args:
  - --domain=cluster.local.
  - --dns-port=10053
  - --config-dir=/kube-dns-config
  - --v=2
  env:
  - name: PROMETHEUS_PORT
    value: "10055"
  ports:
  - containerPort: 10053
    name: dns-local
    protocol: UDP
  - containerPort: 10053
    name: dns-tcp-local
    protocol: TCP
  - containerPort: 10055
    name: metrics
```

9%



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

9%

protocol: TCP
volumeMounts:

```
- name: kube-dns-config
  mountPath: /kube-dns-config
- name: dnsmasq
  image: k8s.gcr.io/k8s-dns-dnsmasq-nanny- (http://k8s.gcr.io/k8s-dns-dns
  livenessProbe:
    httpGet:
      path: /healthcheck/dnsmasq
      port: 10054
      scheme: HTTP
    initialDelaySeconds: 60
    timeoutSeconds: 5
    successThreshold: 1
    failureThreshold: 5
  args:
    - -v=2
    - -logtostderr
    - --configDir=/etc/k8s/dns/dnsmasq-nanny
    - --restartDnsmasq=true
    - --
    - -k
    - --cache-size=1000
    - --no-negcache
    - --log-facility=-
    - --server=/cluster.local/127.0.0.1#10053
    - --server=/in-addr.arpa/127.0.0.1#10053
    - --server=/ip6.arpa/127.0.0.1#10053
  ports:
    - containerPort: 53
    name: dns
```



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

protocol: UDP

containerPort: 53

name: dns-tcp

protocol: TCP

see: <https://github.com/kubernetes/kubernetes/issues/29055> (<https://github.com/kubernetes/kubernetes/issues/29055>)

resources:

requests:

cpu: 150m

memory: 20Mi

volumeMounts:

- name: kube-dns-config

mountPath: /etc/k8s/dns/dnsmasq-nanny

- name: sidecar

image: k8s.gcr.io/k8s-dns-sidecar (<http://k8s.gcr.io/k8s-dns-sidecar>).

livenessProbe:

httpGet:

path: /metrics

port: 10054

scheme: HTTP

initialDelaySeconds: 60

timeoutSeconds: 5

successThreshold: 1

failureThreshold: 5

args:

- --v=2

- --logtostderr

- --probe=kubedns,127.0.0.1:10053,kubernetes.default.svc.cluster.local,

- --probe=dnsmasq,127.0.0.1:53,kubernetes.default.svc.cluster.local,5,S

ports:

- containerPort: 10054

name: metrics

9%



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

protocol: TCP

resources:

requests:

memory: 20Mi

cpu: 10m

dnsPolicy: Default *# Don't use cluster DNS.*

serviceAccountName: kube-dns

9%

2- Deploy kube-dns.

```
$ kubectl create -f kube-dns.yaml
```

3- Verify the deployment.

```
$ kubectl get pods -l k8s-app=kube-dns -n kube-system
```

Installing the Kubernetes dashboard

1- Create the Kubernetes dashboard manifest.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

```
$ vim kube-dashboard.yaml
apiVersion: v1
kind: Secret
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard-certs
  namespace: kube-system
type: Opaque
---
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
---
kind: Role
apiVersion: rbac.authorization.k8s.io/v1 (http://rbac.authorization.k8s.io/v1)
metadata:
  name: kubernetes-dashboard-minimal
  namespace: kube-system
rules:
  # Allow Dashboard to create 'kubernetes-dashboard-key-holder' secret.
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["create"]
  # Allow Dashboard to create 'kubernetes-dashboard-settings' config map.
- apiGroups: [""]
  resources: ["configmaps"]
```

9%



(<http://blog.inkubate.io>)

```
verbs: ["create"]
# Allow Dashboard to get, update and delete Dashboard exclusive secrets.
apiGroups: [""]
resources: ["secrets"]
resourceNames: ["kubernetes-dashboard-key-holder", "kubernetes-dashboard-cert"]
verbs: ["get", "update", "delete"]
# Allow Dashboard to get and update 'kubernetes-dashboard-settings' config map
- apiGroups: [""]
  resources: ["configmaps"]
  resourceNames: ["kubernetes-dashboard-settings"]
  verbs: ["get", "update"]
# Allow Dashboard to get metrics from heapster.
- apiGroups: [""]
  resources: ["services"]
  resourceNames: ["heapster"]
  verbs: ["proxy"]
- apiGroups: [""]
  resources: ["services/proxy"]
  resourceNames: ["heapster", "http:heapster:", "https:heapster:"]
  verbs: ["get"]
---
apiVersion: rbac.authorization.k8s.io/v1 (http://k8s.io/v1)
kind: RoleBinding
metadata:
  name: kubernetes-dashboard-minimal
  namespace: kube-system
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: kubernetes-dashboard-minimal
subjects:
```

9%



(<http://blog.inkubate.io>)

- kind: ServiceAccount
Deploy Kubernetes 1.9 from scratch on VMware vSphere
name: kubernetes-dashboard
0116 words - 23 min read.
namespace: kube-system

9%

```
---
kind: Deployment
apiVersion: apps/v1beta2
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
spec:
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      k8s-app: kubernetes-dashboard
  template:
    metadata:
      labels:
        k8s-app: kubernetes-dashboard
    spec:
      containers:
        - name: kubernetes-dashboard
          image: k8s.gcr.io/kubernetes-dashboard-amd64:v1 (http://k8s.gcr.io/kubernetes-dashboard-amd64:v1)
          ports:
            - containerPort: 8443
              protocol: TCP
          args:
            - --auto-generate-certificates
            # Uncomment the following line to manually specify Kubernetes API ser
```



(<http://blog.inkubate.io>)

```
# If not specified, Dashboard will attempt to auto discover the API s
Deploy Kubernetes 1.9 from scratch on VMware vSphere
# to it. Uncomment only if the default does not work.
# - --apiserver-host=http://my-address:port (http://my-address:port)
volumeMounts:
- name: kubernetes-dashboard-certs
  mountPath: /certs
# Create on-disk volume to store exec logs
- mountPath: /tmp
  name: tmp-volume
livenessProbe:
  httpGet:
    scheme: HTTPS
    path: /
    port: 8443
    initialDelaySeconds: 30
    timeoutSeconds: 30
volumes:
- name: kubernetes-dashboard-certs
  secret:
    secretName: kubernetes-dashboard-certs
- name: tmp-volume
  emptyDir: {}
serviceAccountName: kubernetes-dashboard
# Comment the following tolerations if Dashboard must not be deployed on
tolerations:
- key: node-role.kubernetes.io/master
  effect: NoSchedule
---
kind: Service
apiVersion: v1
metadata:
```

9%



(<http://blog.inkubate.io>)

labels:
Deploy Kubernetes 1.9 from scratch on VMware vSphere
k8s-app: kubernetes-dashboard

6116 words • 23 min read.

```
name: kubernetes-dashboard
namespace: kube-system
spec:
  ports:
    - port: 443
      targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
```

9%

2- Deploy the dashboard.

```
$ kubectl create -f kubernetes-dashboard.yaml
```

Installing Heapster

1- Create a manifest for Heapster.



(<http://blog.inkubate.io>)

\$ vi DeployKubernetes1.9 from scratch on VMware vSphere

apiVersion: v1

kind: ServiceAccount

metadata:

name: heapster

namespace: kube-system

apiVersion: extensions/v1beta1

kind: Deployment

metadata:

name: heapster

namespace: kube-system

spec:

replicas: 1

template:

metadata:

labels:

task: monitoring

k8s-app: heapster

spec:

serviceAccountName: heapster

containers:

- name: heapster

image: k8s.gcr.io/heapster-amd64:v1.4.2 (<http://k8s.gcr.io/heapster-amd64:v1.4.2>)

imagePullPolicy: IfNotPresent

command:

- /heapster

- --source=kubernetes:<https://kubernetes.default> (<https://kubernetes.default>)

apiVersion: v1

kind: Service

9%





(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

```
labels:
  task: monitoring
  # For use as a Cluster add-on (https://github.com/kubernetes/kubernetes/tree
  # If you are NOT using this as an addon, you should comment out this line.
  kubernetes.io/cluster-service (http://kubernetes.io/cluster-service): 'true'
  kubernetes.io/name (http://kubernetes.io/name): Heapster
name: heapster
namespace: kube-system
spec:
  ports:
    - port: 80
      targetPort: 8082
  selector:
    k8s-app: heapster
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1 (http://k8s.io/v1beta1)
metadata:
  name: heapster
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:heapster
subjects:
- kind: ServiceAccount
  name: heapster
  namespace: kube-system
```



(http://blog.inkubate.io)

Access the Kubernetes dashboard

1- Create an admin user manifest.



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

apiVersion: v1

kind: ServiceAccount

metadata:

name: admin-user

namespace: kube-system

apiVersion: rbac.authorization.k8s.io/v1beta1 (<http://rbac.authorization.k8s.io>)

kind: ClusterRoleBinding

metadata:

name: admin-user

roleRef:

apiGroup: rbac.authorization.k8s.io

kind: ClusterRole

name: cluster-admin

subjects:

- kind: ServiceAccount

name: admin-user

namespace: kube-system

2- Create the admin user.

```
$ kubectl create -f kubernetes-dashboard-admin.yaml
```

3- Get the admin user token.

```
$ kubectl -n kube-system describe secret $(kubectl -n kube-system get secret |
```

9%



1- Copy the token Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%

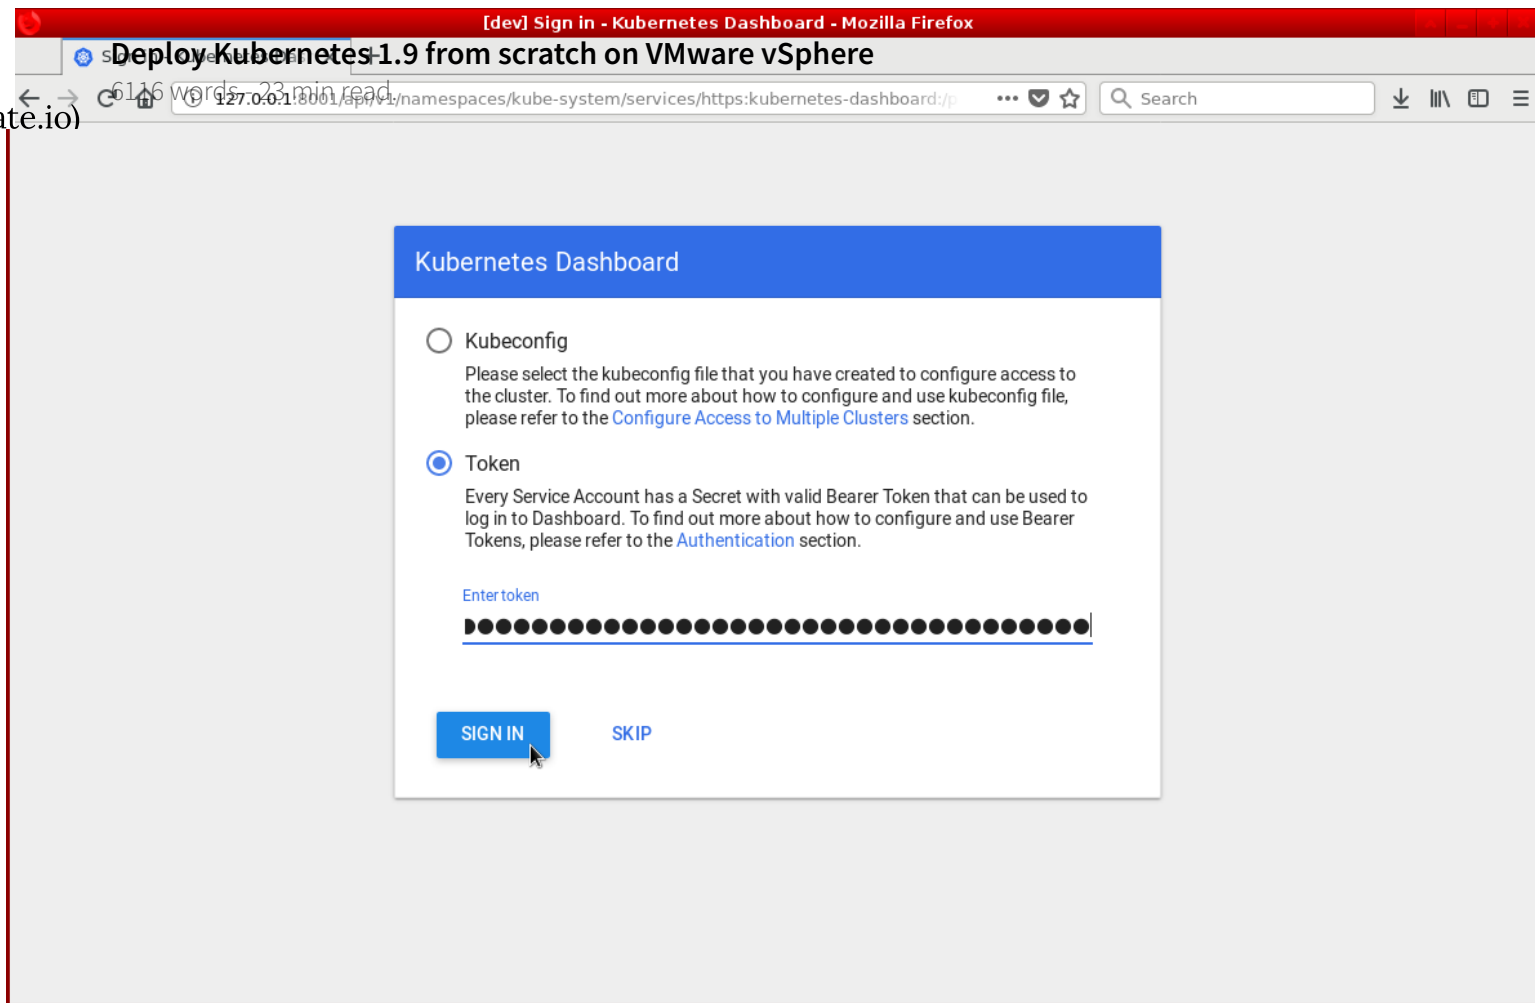
(<http://blog.inkubate.io>)

5- Start the proxy to access the dashboard.

```
$ kubectl proxy
```

6- Browse to <http://localhost:8001/ui> (<http://localhost:8001/ui>).

7- Select Token and paste the token from step 4.



 <http://blog.inkubate.io>

9%


Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read

127.0.0.1:8001/namespaces/kube-system/services/https:kubernetes-dashboard:/p

kubernetes

Search

+ CREATE | 

Overview

Cluster

- Namespaces
- Nodes
- Persistent Volumes
- Roles
- Storage Classes

Namespace

default


Overview

Workloads

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets

Discovery and Load Balancing

Services

Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age
 kubernetes	component: apis. provider: kubern..	10.32.0.1	kubernetes:443 TC kubernetes:0 TCP	-	2 hours


Config and Storage

Secrets

Name	Type	Age
default-token-8784t	kubernetes.io/service-account-token	2 hours

Tags: [Cloud \(/tag/cloud/\)](#), [VMware \(/tag/vmware/\)](#), [vSphere \(/tag/vsphere/\)](#), [Kubernetes \(/tag/kubernetes/\)](#), [Container \(/tag/container/\)](#).

Share this post:

 [https://twitter.com/share?](https://twitter.com/share?text=Deploy%20Kubernetes%201.9%20from%20scratch%20on%20VM)
text=Deploy%20Kubernetes%201.9%20from%20scratch%20on%20VM



kubernetes-1-9-from-scratch-on-vmware-vsphere/)



9%

(http://blog.inkubate.io)

(https://www.facebook.com/sharer/sharer.php?
u=https://blog.inkubate.io/deploy-kubernetes-1-9-from-
scratch-on-vmware-vsphere/)

(https://plus.google.com/share?
url=https://blog.inkubate.io/deploy-kubernetes-1-9-
from-scratch-on-vmware-vsphere/)

« **PREVIOUS** (/DEPLOY-VMWARE-VSPHERE-INTEGRATED-CONTAINERS/)

Deploy VMware vSphere Integrated Containers 1.2.1 (/deploy-vmware-vsphere-integrated-
containers/)

NEXT » (/MONITOR-KUBERNETES-WITH-PROMETHEUS/)

Monitor Kubernetes with Prometheus (/monitor-kubernetes-with-prometheus/)



(<http://blog.inkubate.io>)

Deploy Kubernetes 1.9 from scratch on VMware vSphere

6116 words - 23 min read.

9%



Author

SIMON GUYENNET (/AUTHOR/SIMON/)

Comments

ALSO ON BLOG.INKUBATE.IO

Access your Kubernetes cluster ...

3 years ago • 30 comments

Out of the box, the
Kubernetes authentication is
not very user-friendly for ...

Configure Swift on OpenStack Ocata ...

4 years ago • 3 comments

Swift is an OpenStack
project which provides
object storage to your ...

Automatically generate signed ...

3 years ago • 6 comments


In this lab, we will see how
to automatically generate
signed SSL certificates for ...

Monitor Kubern with Prometheus

4 years ago • 4 comment

Prometheus is a moni
tool originally created
SoundCloud. The proj





9%

Recommend 2

f Share



OR SIGN UP WITH DISQUS 

Vojtech Vavra • 2 years ago

Thank you for your job but I am missing configuration for "containerd ,cri-containerd" unit files.
Thanks !

^